



SOCKET PROGRAMING

IN LINUX - C

Introduction

- The goal of this presentation is to introduce you with concepts of socket programming.
“ How network sockets works? “
- You are expected to explore further by yourself.

Today's Goals:

1

- High-level Overview

2

- Familiarize yourself with socket API

3

- A simple socket connection scheme

4

- Socket management

5

- Multicast Socket

6


- Introducing the programming assignment



High-level Overview

Client – Server

- Clients
 - Locates the server
 - Initiate connection
- Server:
 - Responder
 - Provides service



Client – Server: some key differences

Clients

- ☐ Simple
- ☐ (Usually) sequential
- ☐ Not performance sensitive
- ☐ Execute on-demand

Server

- ☐ Complex
- ☐ (Massively) concurrent
- ☐ High performance
- ☐ Always-on

Client – Server: Similarities

- ❑ Share common protocols
 - Application layer
 - Transport layer
 - Network layer
- ❑ Both rely on APIs for network access

What is API?

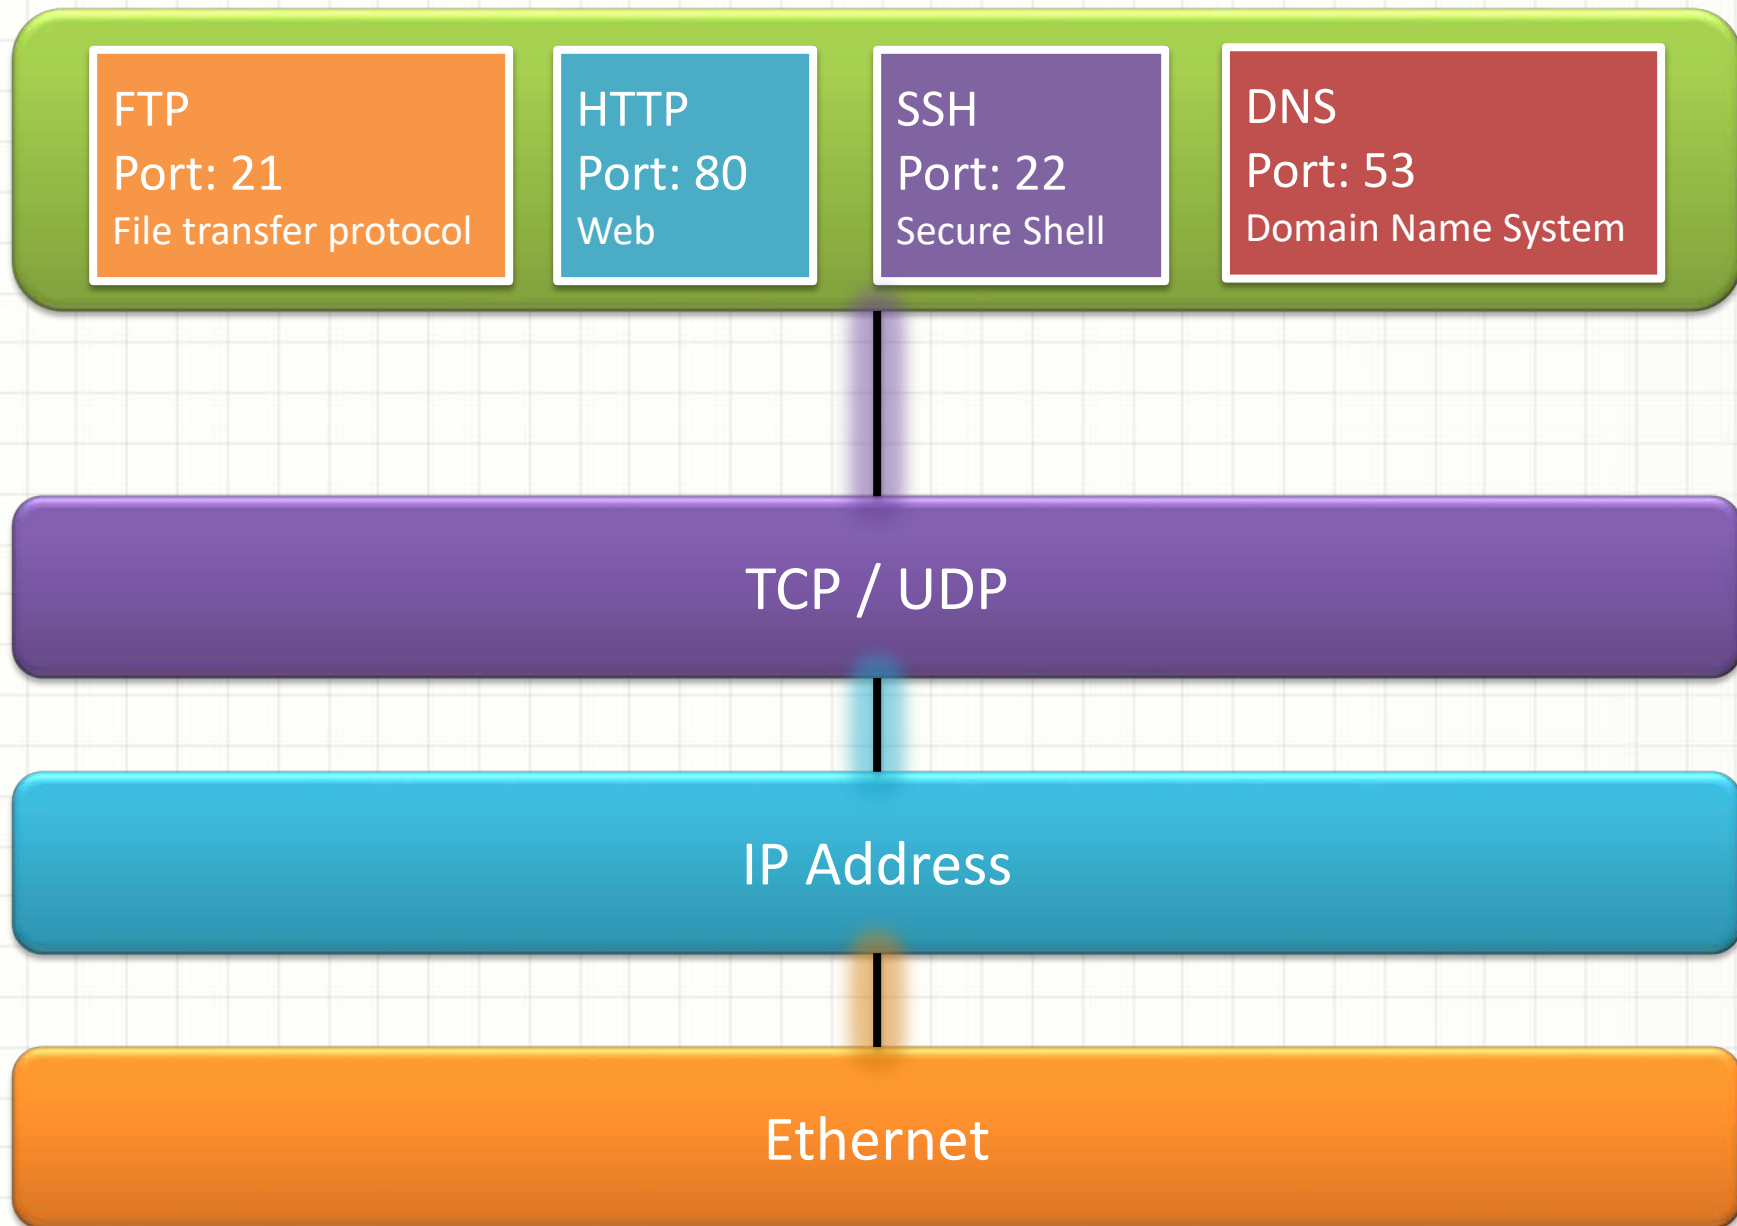
application programming interface (API) is a set of routines, protocols, and tools for building software applications

What is a network socket?

- ☐ It is an application Mailbox for the network messages.
- ☐ Implements an Incoming and Outgoing queues.
- ☐ Managed by the operating system.
- ☐ Represented as file descriptor.
- ☐ Used to pass messages among applications on different computers.
- ☐ Identified by IP address and port.

Addresses, Ports and Sockets

- Like apartments and mailboxes:
 - You are the application
 - Your apartment address is the IP address
 - Your mailbox is the port
 - The post-office is the network
 - Your postman is the operating system
 - The socket is the key that gives you access to the right mailbox

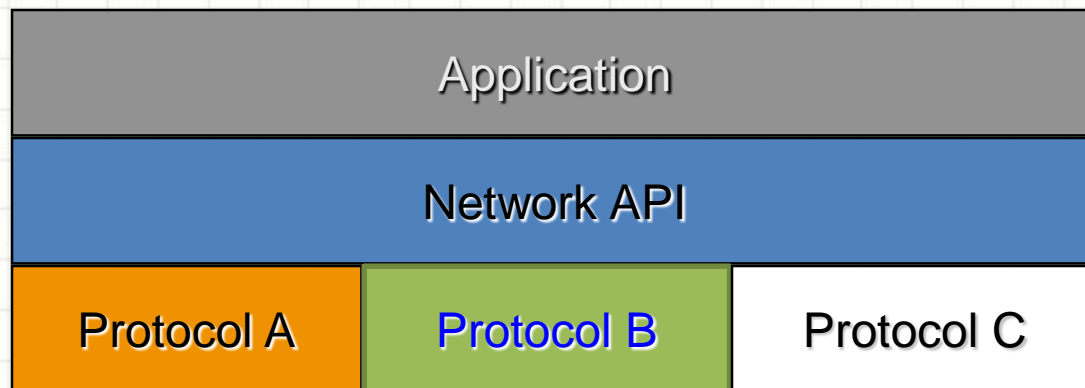




SOCKET API

Introduction

- A **socket API** is an application programming interface (API), provided by the operating system, that allows application programs to control and use network sockets.



Client

Create a socket

Find the server address

Connect to the server

Read/write data

Shutdown connection

Server

Create a socket

Bind the socket

Listen for connections

Accept new client connections

Read/write to client connections

Shutdown connection

socket: creates a socket of a given domain, type, protocol (buy a phone)

<http://linux.die.net/man/2/socket>

TCP: `socket(AF_INET, SOCK_STREAM, 0)`

UDP: `socket(AF_INET, SOCK_DGRAM, 0)`

bind: assigns a name to the socket (get a telephone number)

<http://linux.die.net/man/2/bind>

TCP: `bind(lisen_sock, (struct sockaddr *)&listen_addr, sizeof(listen_addr))`

listen: specifies the number of pending connections that can be queued for a server socket. (call waiting allowance)

<http://linux.die.net/man/2/listen>

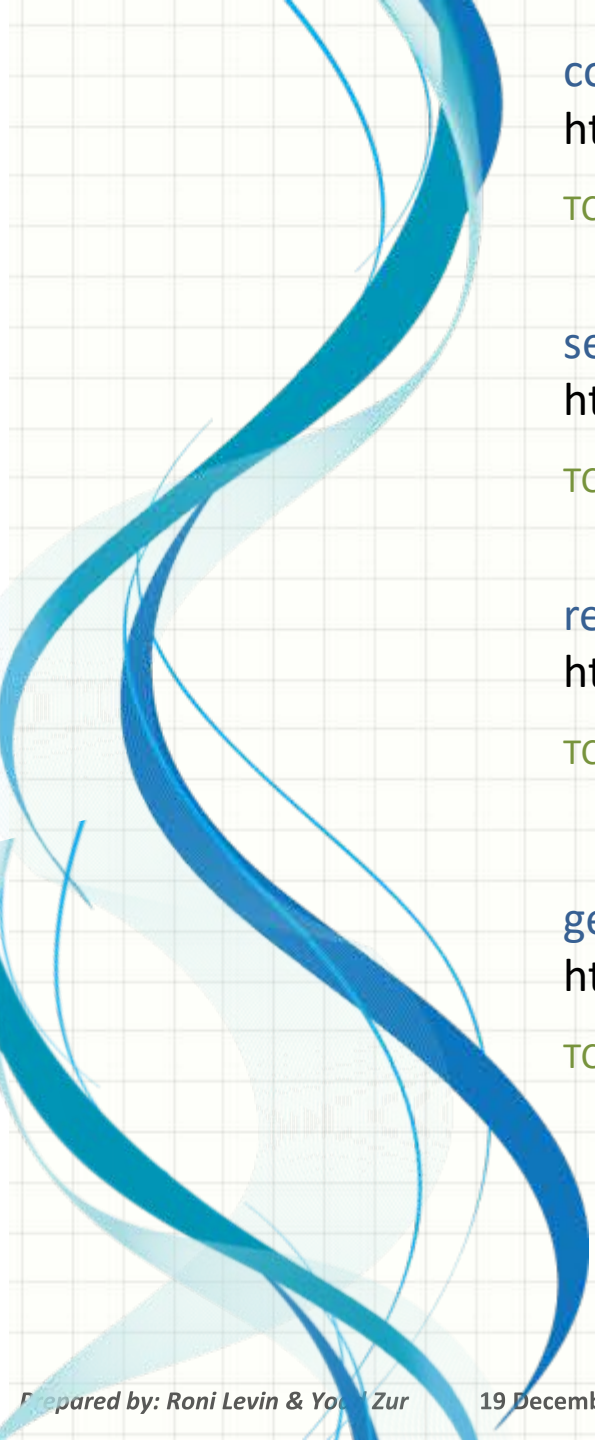
TCP: `listen(lisen_sock, SOMAXCONN)`

accept: server accepts a connection request from a client (answer phone)

<http://linux.die.net/man/2/accept>

TCP: `accept(lisen_sock, (struct sockaddr *)&client_addr, &client_addr_size)`

```
/* The sockaddr structure used for addressing the sockets.
 * It is defined in the file netinet.h.
 */
struct sockaddr {
    uint16_t sin_len; /* length of the structure (16) */
    sa_family_t sin_family; /* AF_INET */
    in_port_t sin_port; /* 16 bit TCP or UDP port number */
    struct in_addr sin_addr; /* 32 bit IPv4 address */
    char sin_zero(8); /* unused */
};
```



connect: client requests a connection request to a server (call)
<http://linux.die.net/man/2/connect>

TCP: `connect(server_sock, (struct sockaddr*) &server_addr, sizeof(server_addr))`

send, sendto: write to connection (speak)
<http://linux.die.net/man/2/send>


TCP/UDP: `send(sock, buffer, &buff_len, 0)`

recv, recvfrom: read from connection (listen)
<http://linux.die.net/man/2/recv>

TCP/UDP: `recv(sock, buffer, &buff_len, 0)`

gethostbyname: find host's address (lookup in phone book)
<http://linux.die.net/man/3/gethostbyname>

TCP/UDP: `struct * = gethostbyname(char *)`



`htons()`, `htonl()`, `ntohs()`, `ntohl()`: convert to/from network address
<http://linux.die.net/man/3/htons>

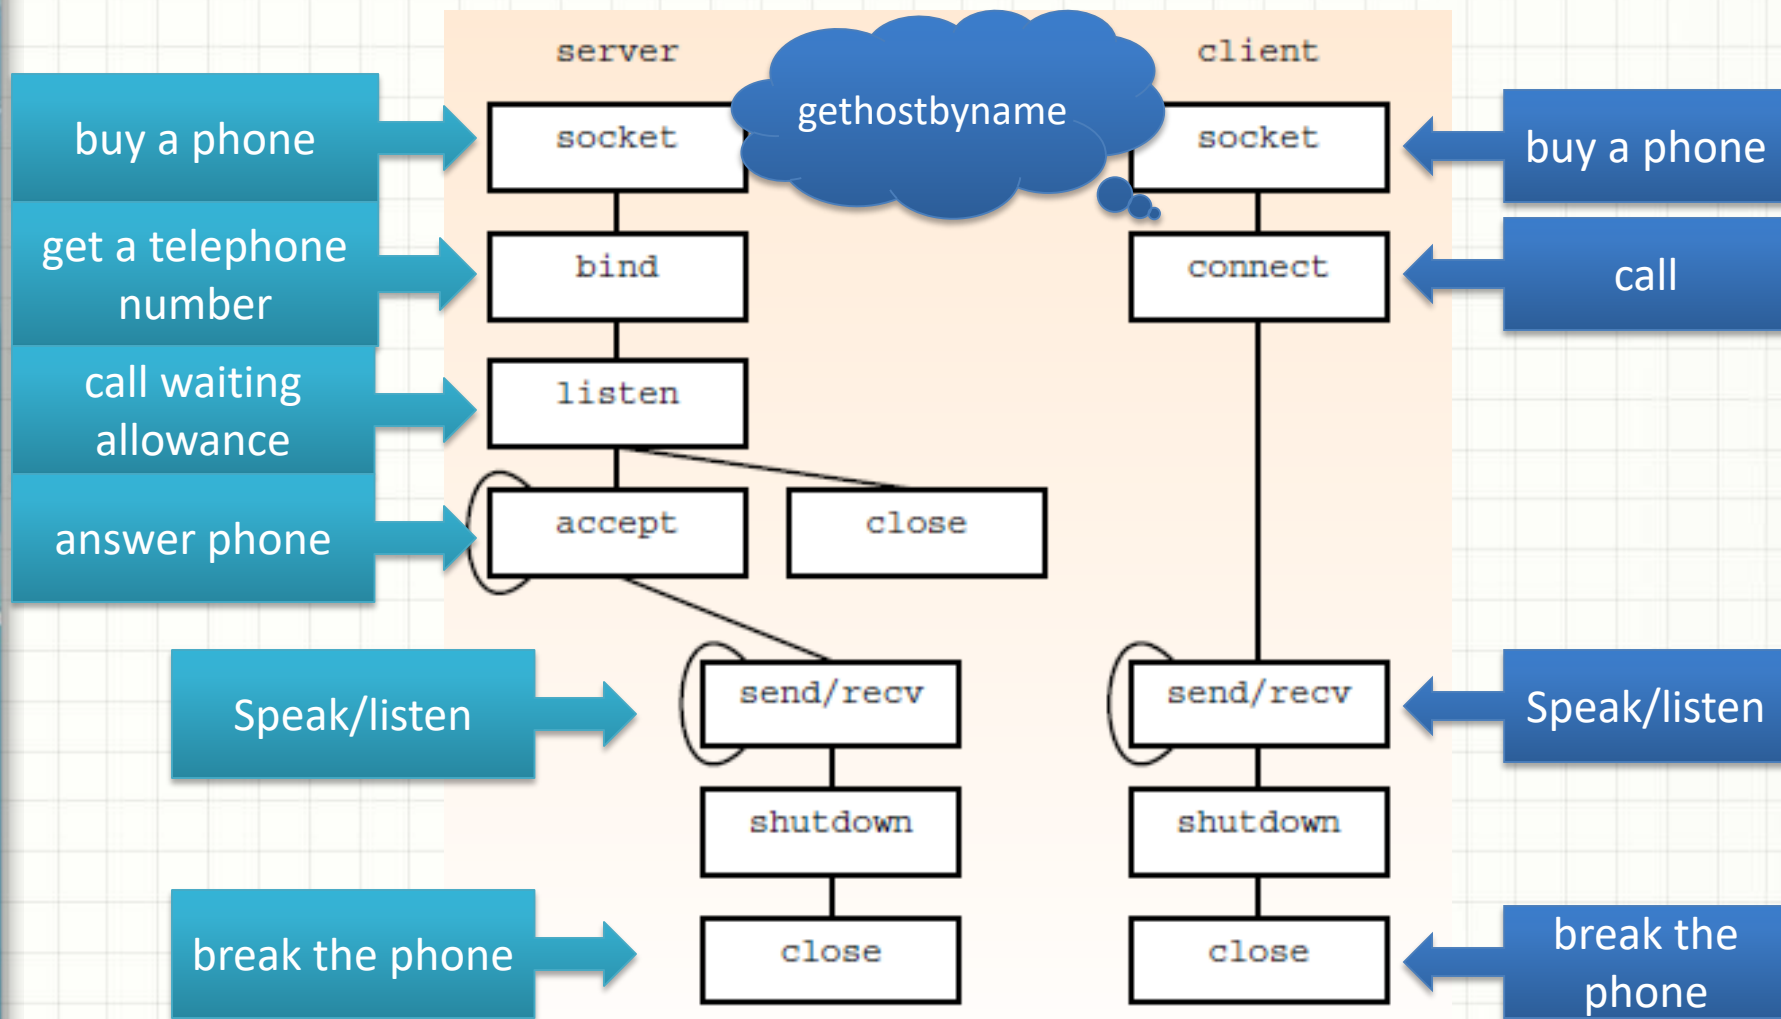
`close`: close the socket (break the phone)
<http://linux.die.net/man/2/close>

TCP/UDP: `close(socket)`

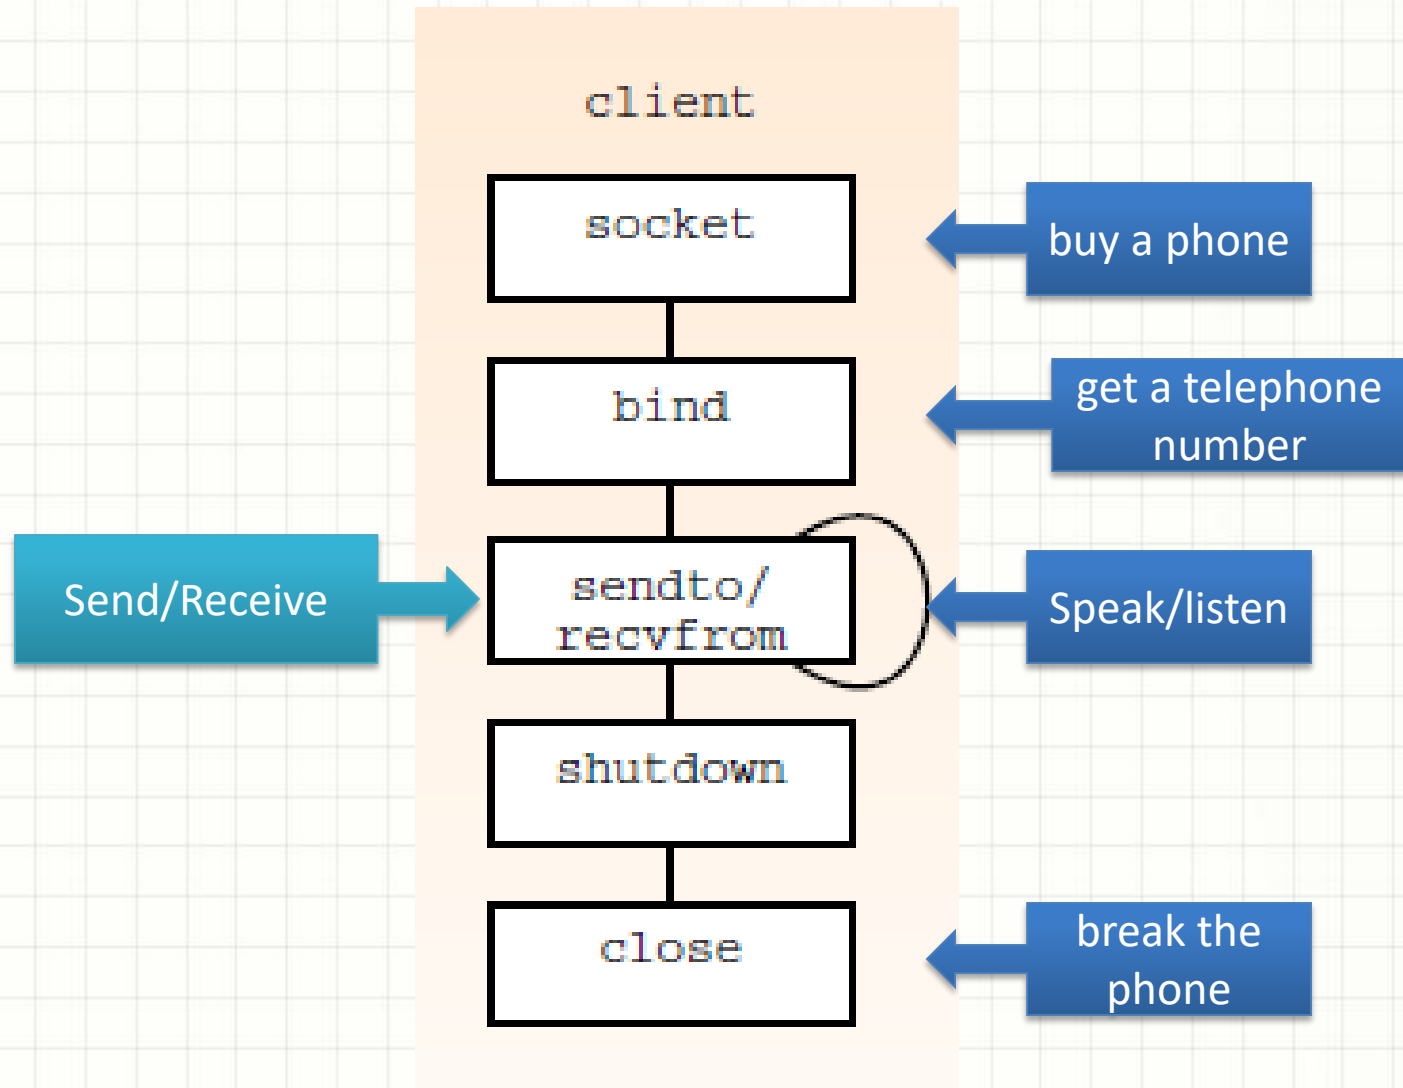


CONNECTION SCHEME

TCP



UDP





SOCKET MANAGEMENT

Socket management methods

- Parallel: Threads / process
- Serial: select function

select: waiting for sockets change

<http://linux.die.net/man/2/select>


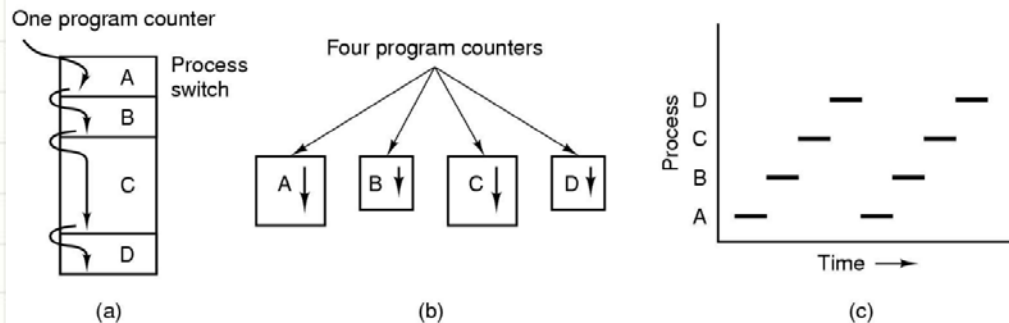
TCP/UDP: `int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);`

We define a set of sockets which in turn will be handled.

```
void FD_CLR(int fd, fd_set *set);  
int FD_ISSET(int fd, fd_set *set);  
void FD_SET(int fd, fd_set *set);  
void FD_ZERO(fd_set *set);
```

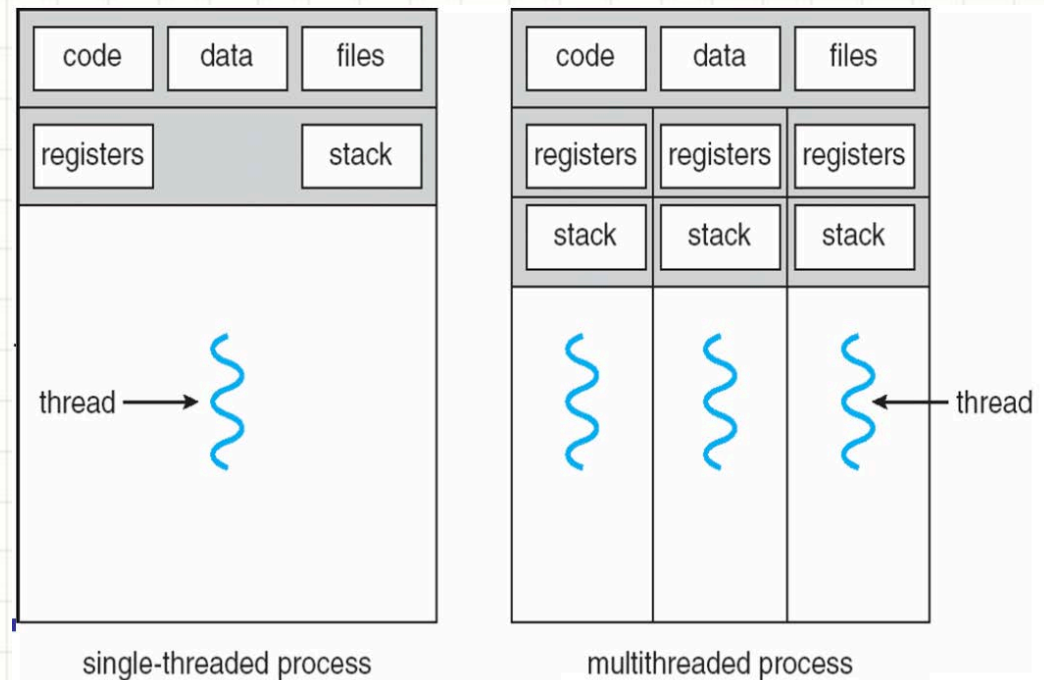
Processes and Threads

- A process is a program that runs sequentially
- When we have several processes



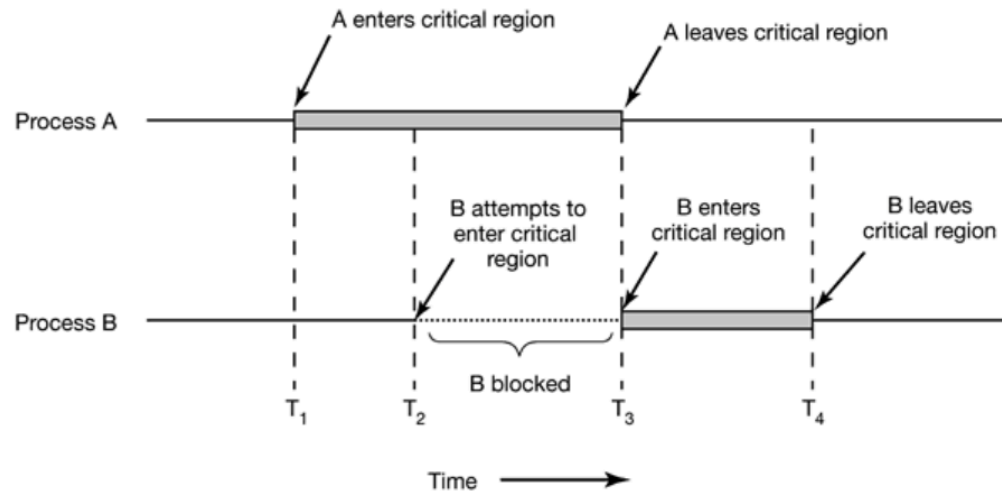
- Each process has
 - address space – code, variables and etc.
 - Single thread of control

- Multithreading
 - it is desirable to have multiple threads of control
 - in the same address space
 - running in quasi-parallel



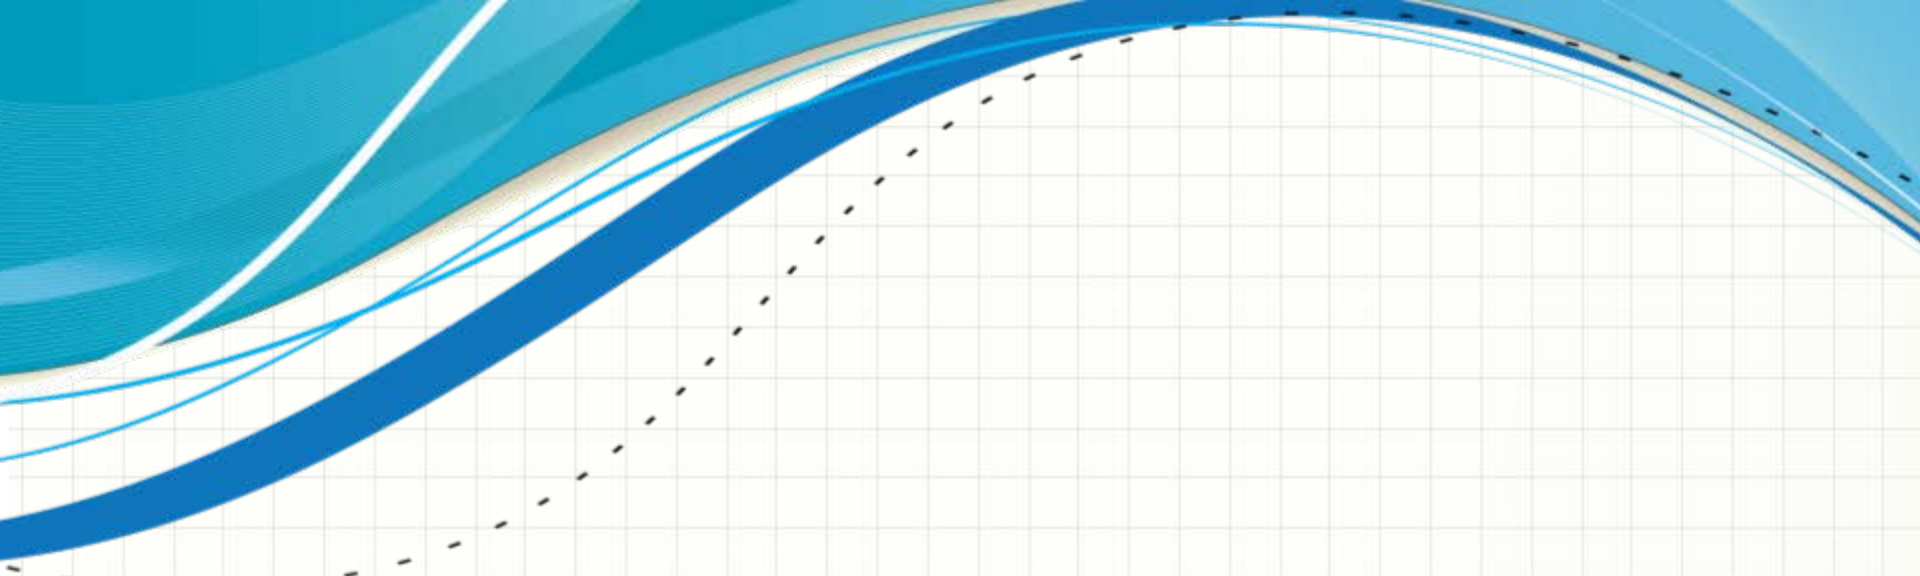
- No need in context switch

Mutual Exclusion using Critical Regions



Mutual exclusion using critical regions

- Mutex: a variable that can be in one of two states
 - Unlocked
 - locked



MULTICAST SOCKET

Using Multicast capabilities

- Multicast source – Only need to send multicast packets with multicast address as destination

```
sock = socket(AF_INET, SOCK_DGRAM, 0);  
addr.sin_family = AF_INET;  
addr.sin_addr.s_addr = inet_addr("239.0.0.1");  
addr.sin_port = htons(6000);  
sendto(sock, message, sizeof(message), 0, (struct sockaddr *) &addr, sizeof(addr));
```

- The network does all the work (PIM)

Using Multicast capabilities

- Multicast receiver – need to subscribe to the multicast group (tell the kernel which multicast groups you are interested in)

```
struct ip_mreq mreq;
```

```
sock = socket(AF_INET, SOCK_DGRAM, 0);  
addr.sin_family = AF_INET;  
addr.sin_addr.s_addr = htonl(INADDR_ANY);  
addr.sin_port = htons(6000);
```

```
bind(sock, (struct sockaddr *) &addr, sizeof(addr))  
mreq.imr_multiaddr.s_addr = inet_addr("239.0.0.1");  
mreq.imr_interface.s_addr = htonl(INADDR_ANY);  
setsockopt(sock, IPPROTO_IP, IP_ADD_MEMBERSHIP, &mreq, sizeof(mreq))
```

IP_DROP_MEMBERSHIP

```
recvfrom(sock, message, sizeof(message), 0, (struct sockaddr *) &addr, &addrlen);
```

- The network does all the work (IGMP, PIM)



FINAL PROJECT:

INTERNET RADIO APPLICATION

Radio over Multicast

- Implemented over the multicast lab
- Server
 - Multicast stations: each station a different multicast group, all use same port (UDP)
 - Sends info to clients (TCP)
- Two Clients
 - Listener: a simple UDP and TCP receiver
 - Controller: talk to server to find info (using TCP) and inform Listener (using TCP).

Internet Radio Application

- You will implement 3 programs, step by step:
 1. Client: Listener
 2. Client: Controller
 3. Server
- We will provide binaries of all 3 of them so you can check your programs one by one

Client: Listener

- UDP socket
- Receives audio packets from the server
- TCP socket
- Receives info commands from the Client Controller

Client: Controller

- Two TCP sockets
- User input
- The user interface with the server
- Handshake with server to learn the multicast address
- Can ask about songs
- Can ask for listeners list
- Inform Listener for the proper song


Server

- TCP and UDP socket
- Maintains control connections with listeners
- Transmits stations via UDP to multicast
- Maintains several stations
- Important: Rate control


Server – Client Connection Protocol

Client Controller

Hello




AskSong




Server

Welcome




Announce



Server – Client Connection Protocol

Client Controller


Server



Client – Client Connection Protocol

Client Controller

Client Listener



Thanks to...

- Wikipedia
- http://www.ccs.neu.edu/home/cbw/4700/slides/4_C_Sockets.pptx
- http://www.cs.northwestern.edu/~agupta/cs340/sockets/sockets_intro.ppt
- <http://parsys.eecs.uic.edu/~solworth/sockets.pdf>